

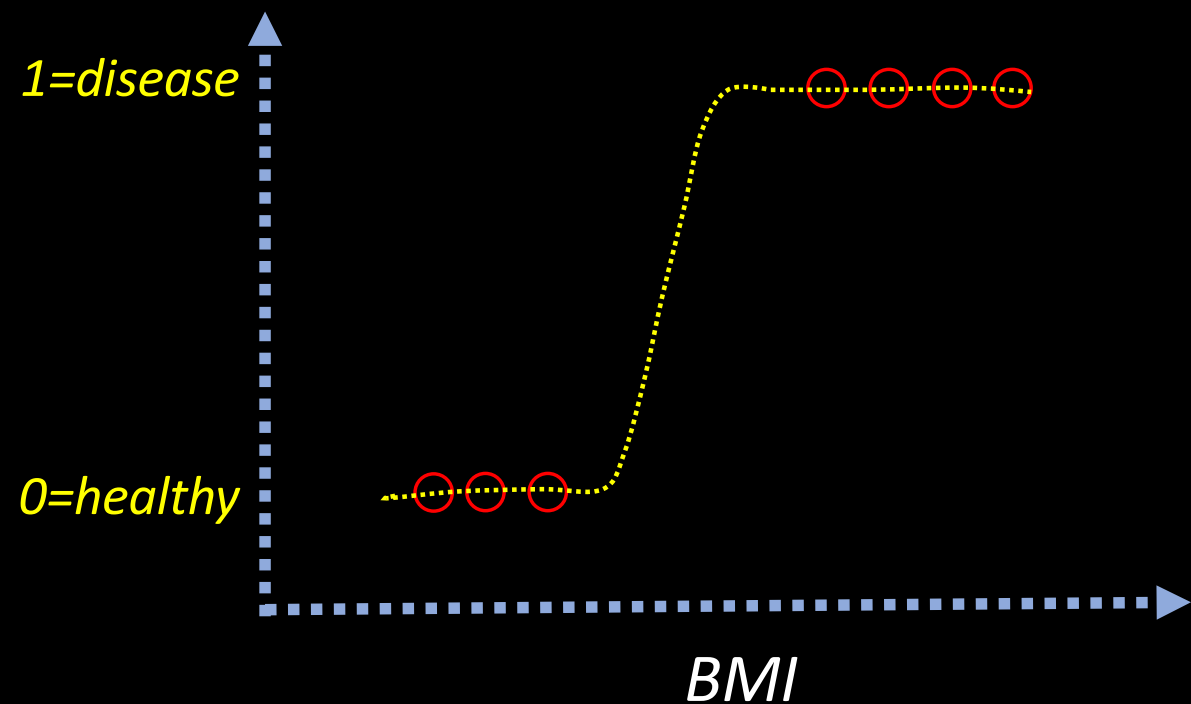
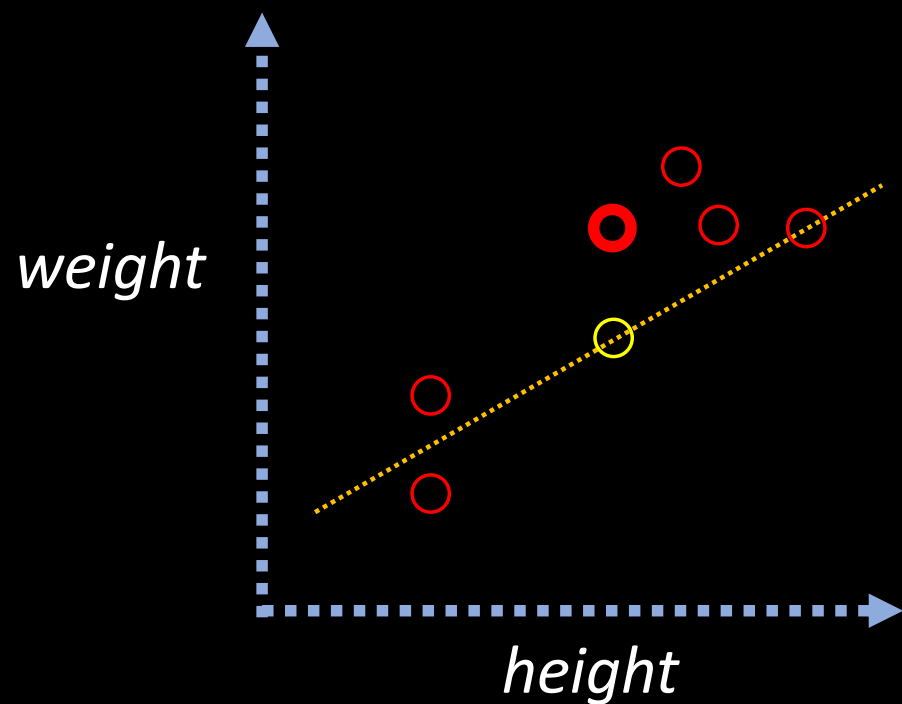
Linear & Logistic Regression

Regression vs Classification

Regression: Predicts continuous outcomes by finding the best-fit line through data points. Used for forecasting and understanding relationships between variables.

Classification: Predicts the probability of an instance belonging to a particular class.

Regression vs Classification



Regression:

Predicting drug dosage based on patient weight

Estimating bone density changes over time in osteoporosis patients

Analyzing the relationship between blood pressure and salt intake

Classification:

Predicting the likelihood of heart disease based on cholesterol levels

Determining the probability of remission in cancer patients given certain treatment

Classifying medical images as showing presence or absence of a tumor

Linear Regression: Predicting drug dosage based on patient weight

```
from sklearn.linear_model import LinearRegression  
import numpy as np
```

Sample data: patient weights (kg) and corresponding drug dosages (mg)

```
weights = np.array([[60], [70], [80], [90], [100]])  
dosages = np.array([100, 115, 130, 145, 160])
```

Code Block 1

```
model = LinearRegression()
```

```
model.fit(weights, dosages)
```

```
# Predict dosage for a new patient weighing 75 kg
```

```
new_weight = np.array([[75]])
```

```
predicted_dosage = model.predict(new_weight)
```

```
print(f"Predicted dosage for 75 kg: {predicted_dosage[0]:.2f} mg")
```

```
# Output: ``Predicted dosage for 75 kg: 122.50 mg ``
```

Code Block 1

Logistic Regression: Predicting presence of heart disease

```
from sklearn.linear_model import LogisticRegression
```

Sample data: patient cholesterol levels and heart disease presence (0: No, 1: Yes)

```
cholesterol = np.array([[150], [200], [250], [300], [350]])
```

```
heart_disease = np.array([0, 0, 1, 1, 1])
```

```
log_model = LogisticRegression()
```

```
log_model.fit(cholesterol, heart_disease)
```

Code Block 1

```
# Predict heart disease for a new patient with cholesterol level  
275
```

```
new_cholesterol = np.array([[275]])
```

```
predicted_prob =
```

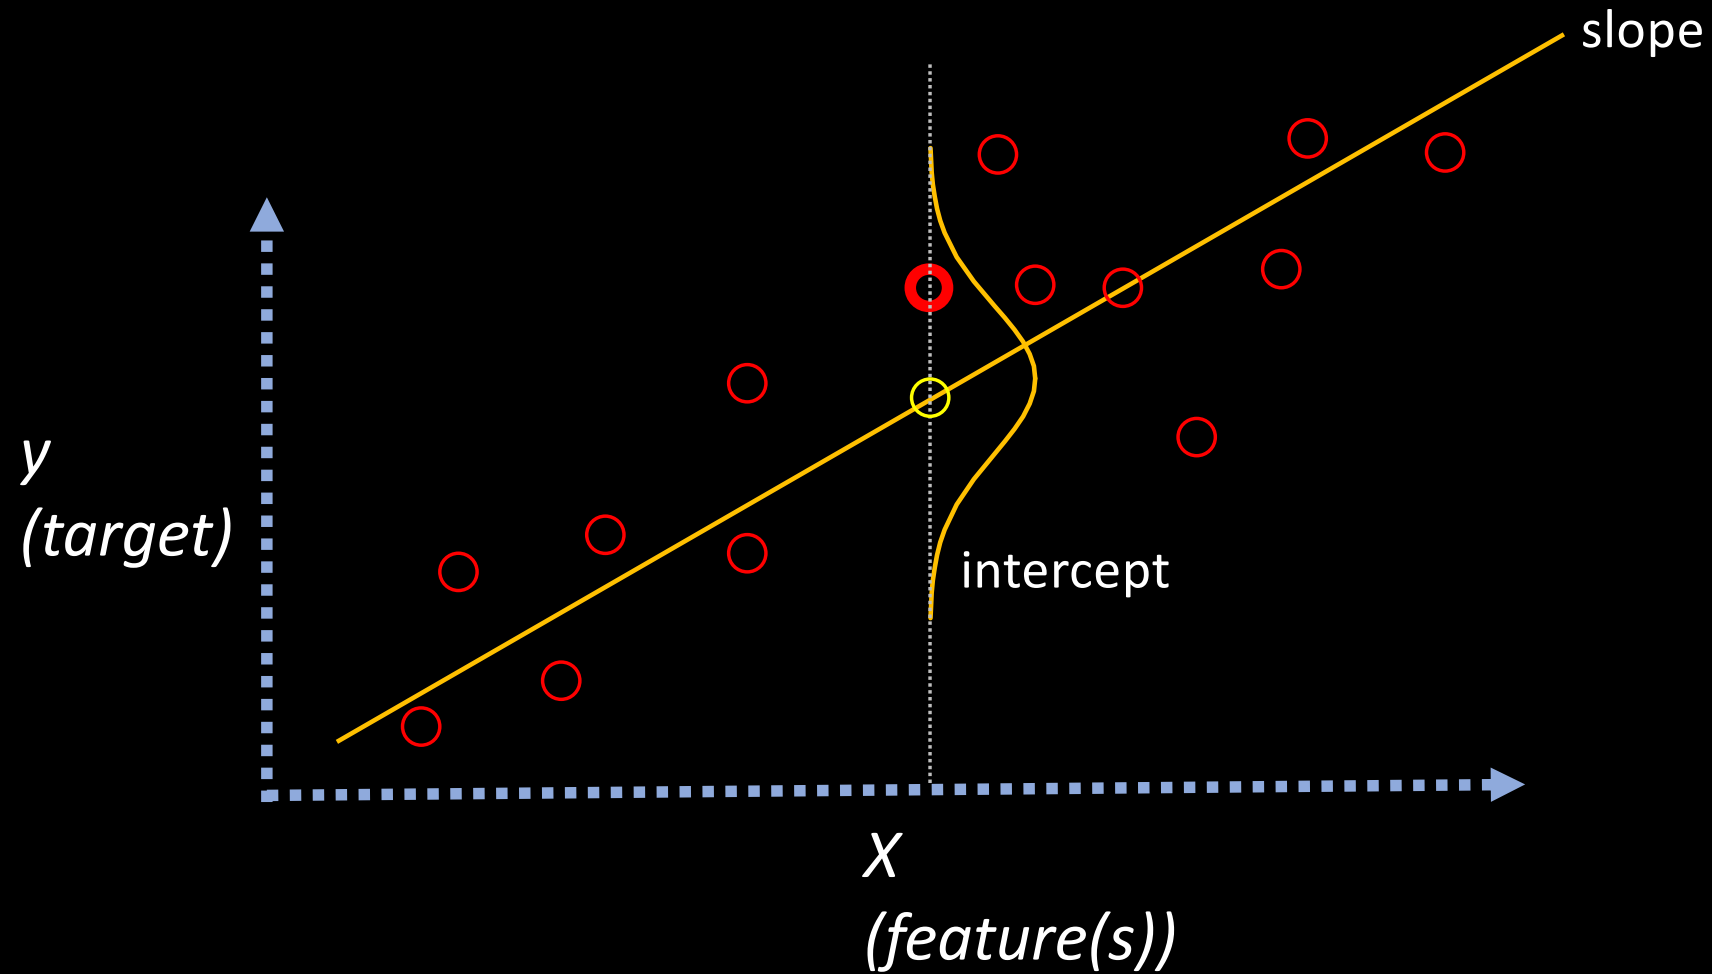
```
log_model.predict_proba(new_cholesterol)[0][1]
```

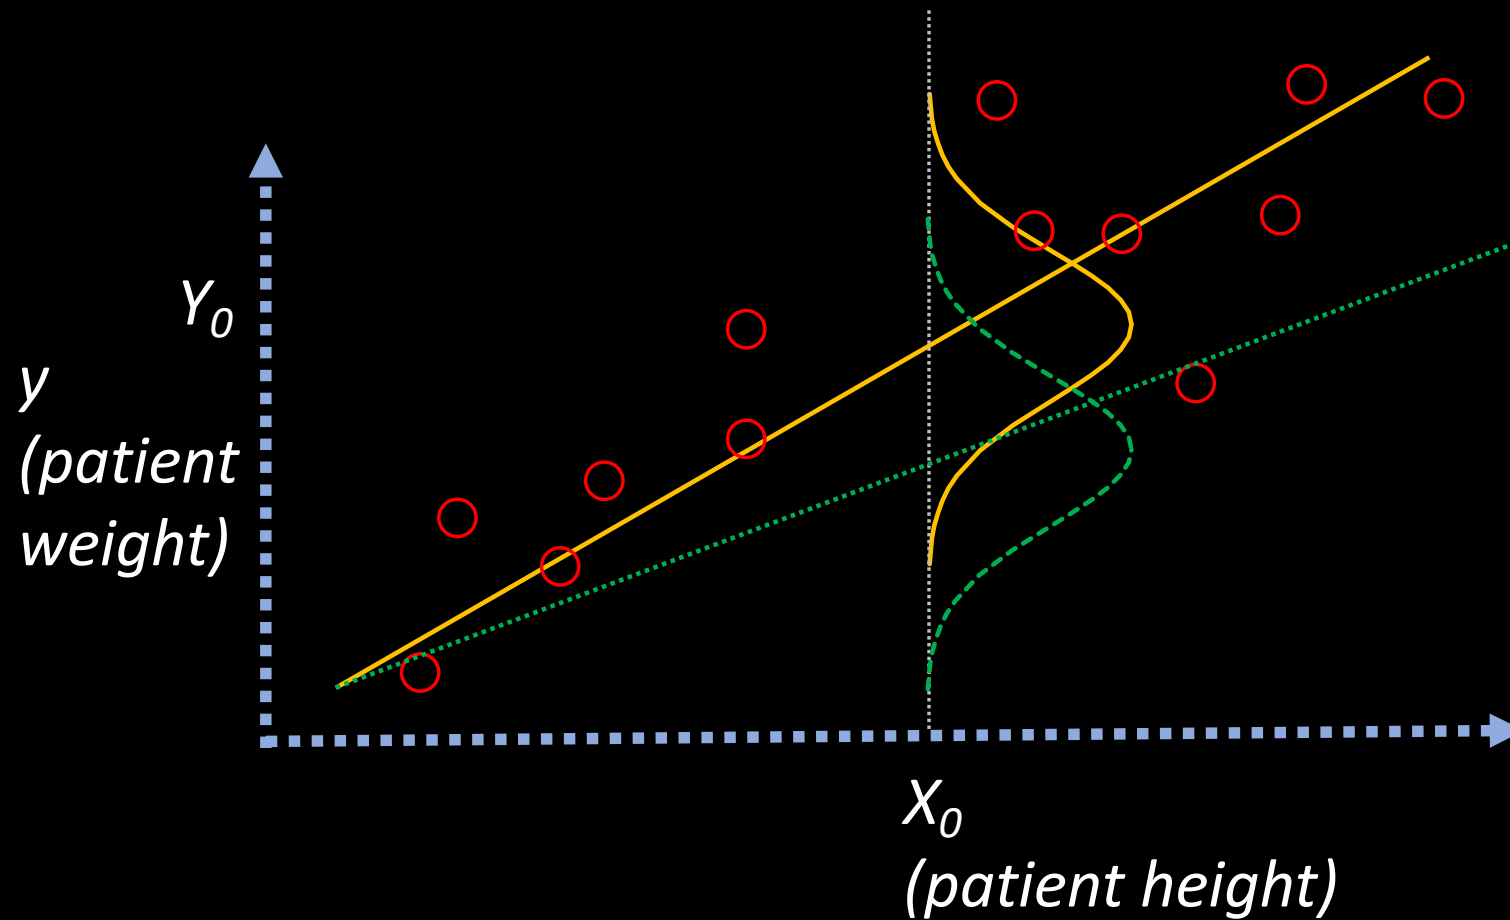
```
print(f"Probability of heart disease for 275 cholesterol:  
{predicted_prob:.2f}")
```

```
# Output: ``Probability of heart disease for 275 cholesterol: 0.73  
``
```

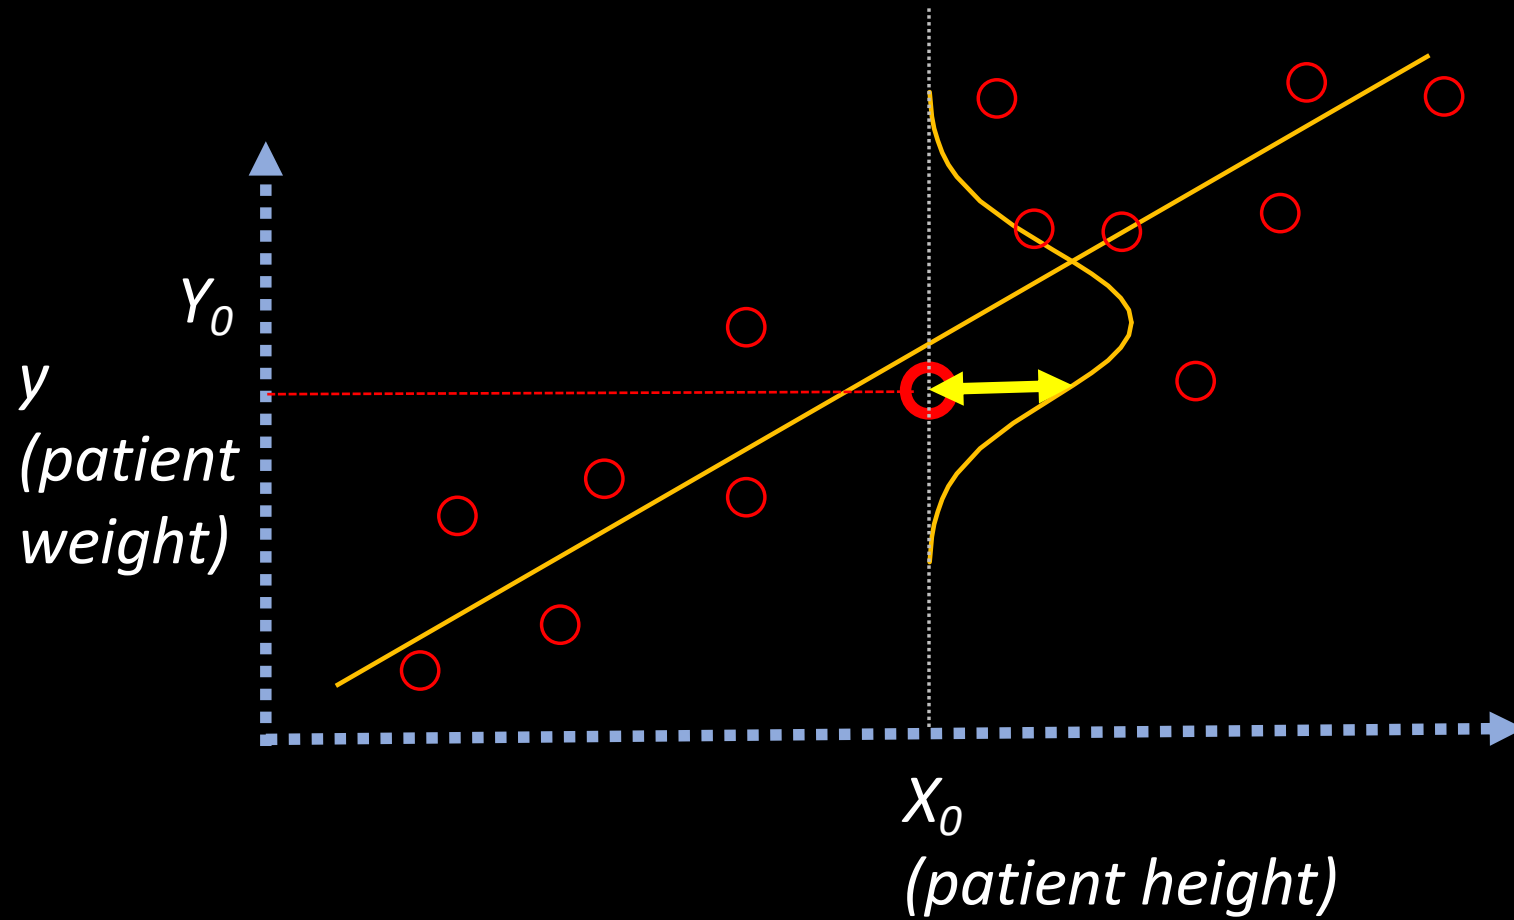
Code Block 1

Linear Regression: Predicts continuous outcomes

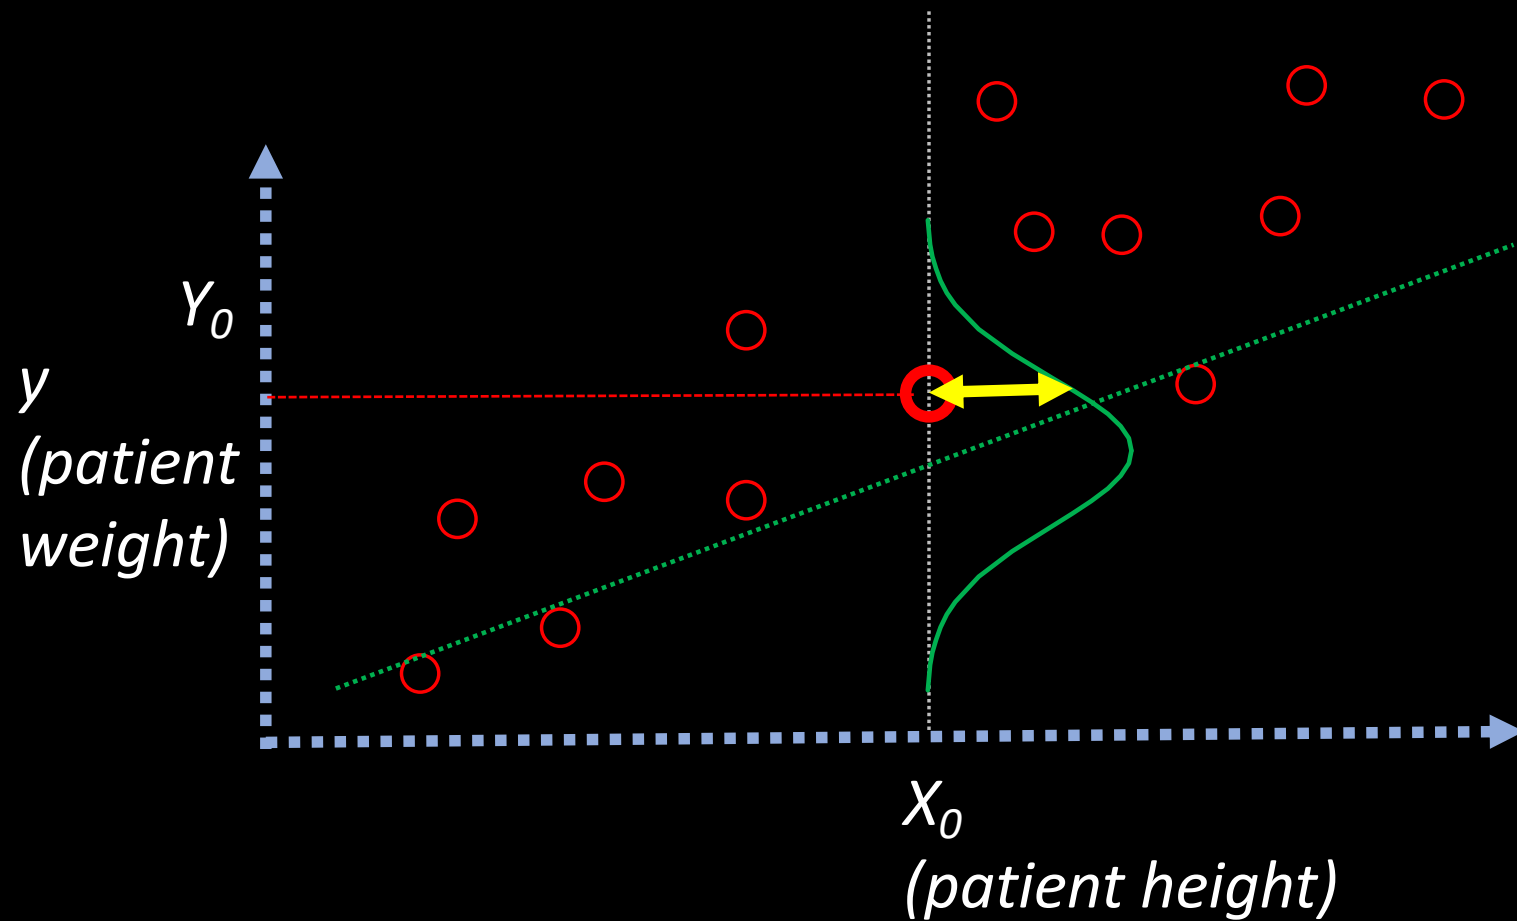




Which model
Is more probable?



Picking the
more probable model



Picking the
more probable model

Maximum Likelihood Estimation (MLE)

$$\theta_{ML} = \operatorname{argmax}_{\theta} P(y|X; \theta) = \operatorname{argmin}_{\theta} (-P(y|X; \theta))$$

MLE is a method for estimating the parameters of a statistical model by maximizing the likelihood function.

Key Concept: It finds the parameter values that make the observed data **most probable**.

Review probability

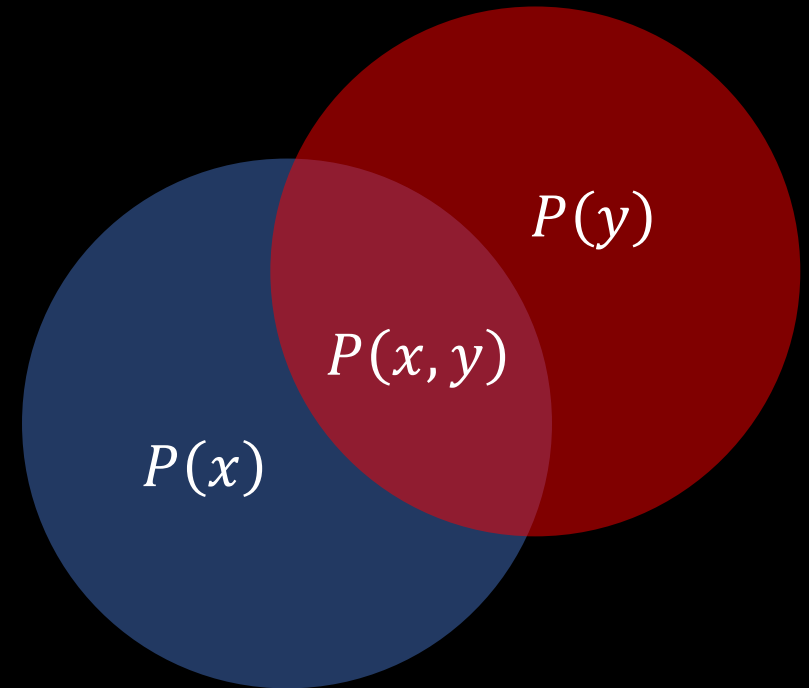
$P(x)$

Probability

$P(x, y)$

Joint Probability

$P(y|x) = P(x, y) / P(x)$ Conditional Probability



Probability:

The likelihood of an event occurring. Example: The probability of a patient responding to a specific drug treatment.

Joint Probability:

The probability of two or more events occurring together. Example: The probability of a patient having both diabetes and hypertension.

Conditional Probability:

The probability of an event occurring, given that another event has already occurred. Example: The probability of a patient having a heart attack, given that they have high blood pressure.

Maximum Likelihood Estimation (MLE)

$$\theta_{MLE} = \operatorname{argmax}_{\theta} P(y|X; \theta)$$

$P(y|X)$ Probability of seeing y given X (true probabilistic distribution)

$P(y|X; \theta)$ Probability of seeing y predicted X from model with parameters θ

$\operatorname{argmax}_{\theta}$ Find the parameters θ to maximize the given probability $P(y|X, \theta)$

Probability

$$P(A|B) \quad :=$$

Chance of seeing data A
given fixed condition B

Likelihood

$$L(B|A)$$

Chance of condition B
given seeing data A

Maximum Likelihood Estimation (MLE)

$$\theta_{MLE} = \operatorname{argmax}_{\theta} P(y|X; \theta)$$

$P(y|X)$ (true) probability of seeing stroke given blood pressure

$P(y|X; \theta)$ Probability of seeing stroke given blood pressure from model with θ

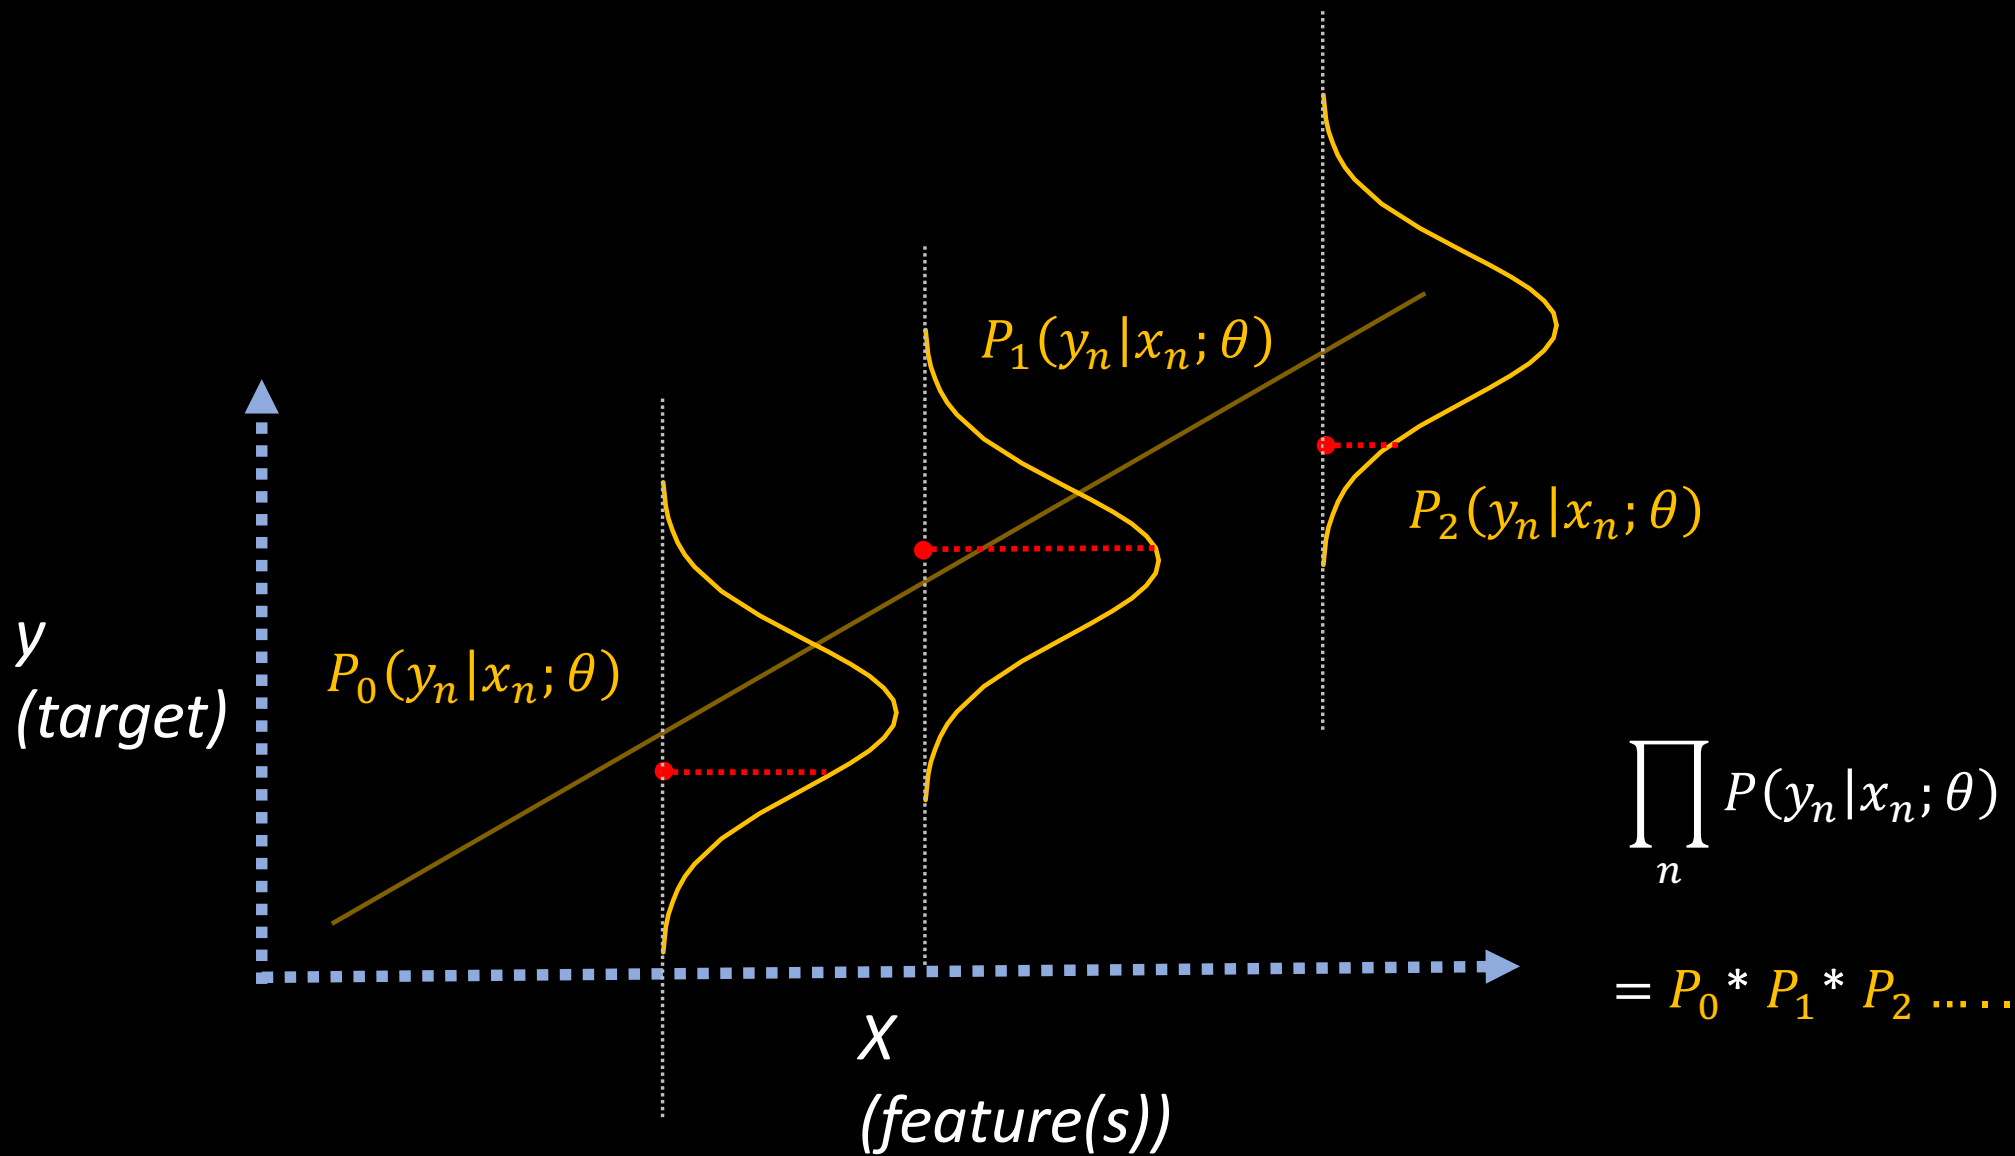
$\operatorname{argmax}_{\theta}$ Find the parameters θ to maximize the given probability $P(y|X, \theta)$

MLE in Linear Regression

Assumption: In linear regression, we assume errors are normally distributed around the true regression line.

Likelihood Function: Measures the probability of observing our data given a particular set of regression parameters (slope and intercept).

MLE Process: Finds the regression line that maximizes this likelihood function, making our observed data most probable, equivalent to minimizing the sum of squared errors (least squares method).



$$P(A, B) = P(A) * P(B)$$

If A and B are independent
(often assumed that data are **independent and identically distributed (IID)**)

$$\text{For } y = \{y_1, y_2 \dots y_i\}, X = \{x_1, x_2 \dots x_n\}$$

Features and targets
for all the
observations

$$P(y|X; \theta) = \text{Likelihood of all the observations}$$

$$P(y_1|x_1; \theta) * P(y_2|x_2; \theta) \dots * P(y_n|x_n; \theta) = \prod_n P(y_n|x_n; \theta)$$

Data points

```
weights = np.array([[60], [70], [80], [90], [100]])
```

```
dosages = np.array([100, 115, 130, 145, 160])
```

Create and train the good model

```
model_good = LinearRegression()
```

```
model_good.fit(weights, dosages)
```

Create a bad model with random coefficients

```
model_bad = LinearRegression()
```

```
model_bad.coef_ = np.array([1.53])
```

```
model_bad.intercept_ = np.array([9.5])
```

Code Block 2

```
def calculate_probabilities(predictions, actual, std_dev):  
    return norm.pdf(actual - predictions, loc=0, scale=std_dev)
```

```
# Predict dosages
```

```
predictions_good = model_good.predict(weights)  
predictions_bad = model_bad.predict(weights)
```

```
probs_good = calculate_probabilities(predictions_good, dosages, 1)  
probs_bad = calculate_probabilities(predictions_bad, dosages, 1)
```

```
print(probs_good) print(probs_bad)
```

Code Block 2

prob of the bad model:

[0.17136859 0.11092083 0.06561581 0.03547459 0.0175283]

prob of the good model:

[0.39894228 0.39894228 0.39894228 0.39894228 0.39894228]

Code Block 2

Joint Probability

$$P(y_1|x_1; \theta) * P(y_2|x_2; \theta) \dots * P(y_n|x_n; \theta) = \prod_n P(y_n|x_n; \theta)$$

$$\prod_n P(y_n|x_n; \theta) \quad (0.1 * 0.1 * 0.1 * 0.1 \dots) \quad 0.0000000001$$

VS

$$\sum_i \log P(y_n|x_n; \theta) \quad (-1 -1 -1 -1 -1 \dots) \quad -9$$

Log Probability

```
log_probs_good = np.log(probs_good)
log_probs_bad = np.log(probs_bad)
print("\nLog probabilities for good model:")
print(log_probs_good)print(log_probs_bad)

# Calculate sum of log probabilities (log-likelihood)
log_likelihood_good = np.sum(log_probs_good)
log_likelihood_bad = np.sum(log_probs_bad)
print(f"\nLog-likelihood for good model: {log_likelihood_good:.2f}")
print(f"Log-likelihood for bad model: {log_likelihood_bad:.2f}")
```

Code Block 3

Log probabilities for good model:

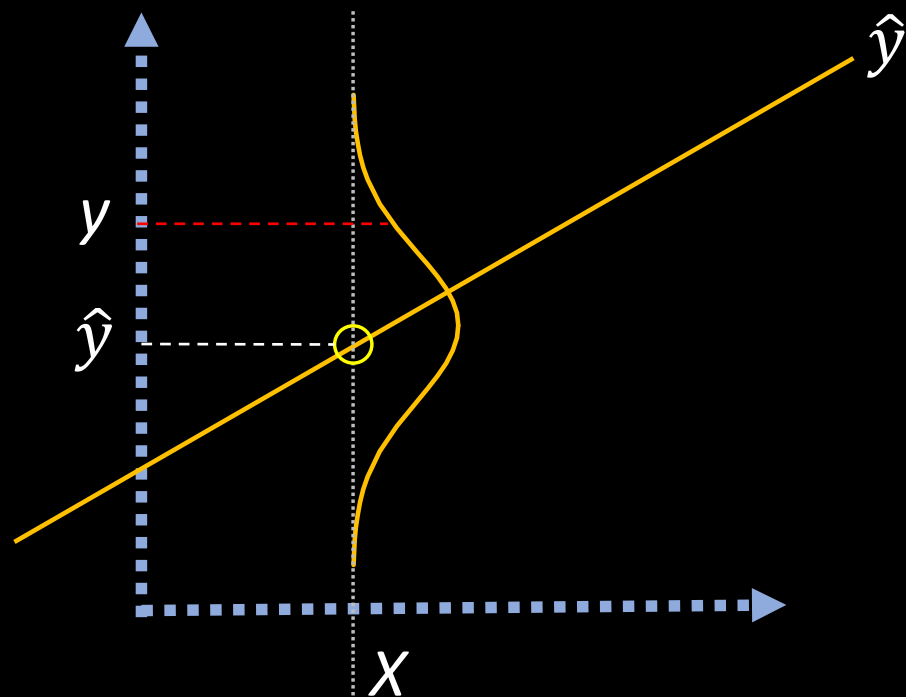
$[-0.91893853 \ -0.91893853 \ -0.91893853 \ -0.91893853 \ -0.91893853]$

Log probabilities for bad model:

$[-1.76393853 \ -2.19893853 \ -2.72393853 \ -3.33893853 \ -4.04393853]$

Log-likelihood for good model: -4.59

Log-likelihood for bad model: -14.07



Picking the most probable parameters (w_i and b) for the linear regression mode, or

the model which has the most **LIKELIHOOD** based on the present observation.

aka Maximum Likelihood Estimation (MLE)

$$P(y|x) = N(\hat{y}, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{y-\hat{y}}{\sigma}\right)^2} \quad \sigma \text{ (variance)}$$

$$\hat{y} = wx + b$$

$$\text{or } = w_1x_1 + w_2x_2 + \dots + w_nx_n + b = \sum_i w_ix_i + b \text{ for multi-variables}$$

$$\sum_n \log P(y_n | x_n; \theta)$$

$$= n \sum \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{(y_n - wx_n - b)^2}{2\sigma^2} \right) \right)$$

Linear Regression picked from Normal Distribution

$$= -n * \log \frac{1}{\sqrt{2\pi\sigma^2}} - \sum_i \frac{(y_n - wx_n - b)^2}{2\sigma^2}$$

Assuming constant variance

$$\operatorname{argmax}(\sum_i \log P(y_n | x_n; \theta)) = \operatorname{argmin}(\sum_n (y_n - wx_n - b)^2)$$

$$= \operatorname{argmin}(\sum_n (y_n - \hat{y})^2)$$

Mean Square Error (MSE)

Linear Regression is optimized by
minimize mean square error (MSE)

Linear Regression and Maximum Likelihood Estimation (MLE)

Keys:

- Likelihood: Probability of observing the data given the model parameters.
- Log-Likelihood: Sum of log probabilities, used for numerical stability and easier optimization.
- Assumptions: Errors are normally distributed around the true regression line.

MLE Process:

- Define the likelihood function based on the probability distribution of errors.
- Find the parameters that maximize the log-likelihood.

Advantages of MLE:

- Provides a principled way to estimate model parameters.
- Allows for comparison between different models using log-likelihoods

Regression by just called a library (Scipy)

```
from sklearn.linear_model import Regression
```

```
Regr = Regression()
```

```
Regr.fit(x_train, y_train)
```

Behind the Scene `Regr.fit(x_train, y_train)`

Ability for the model to
improve itself (optimization)

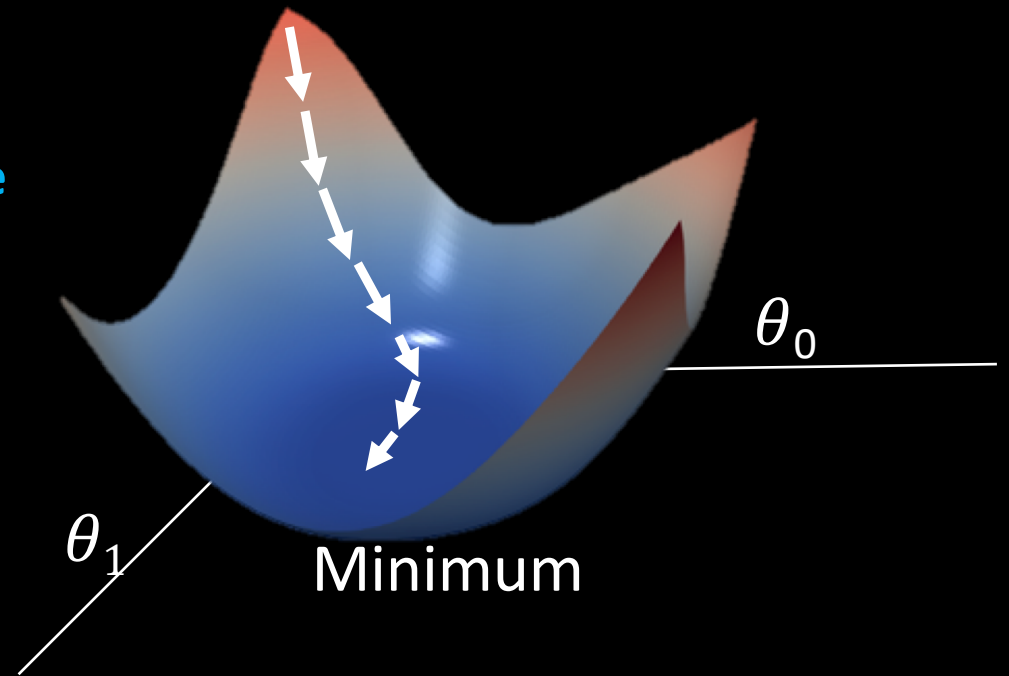


Gradient Descent

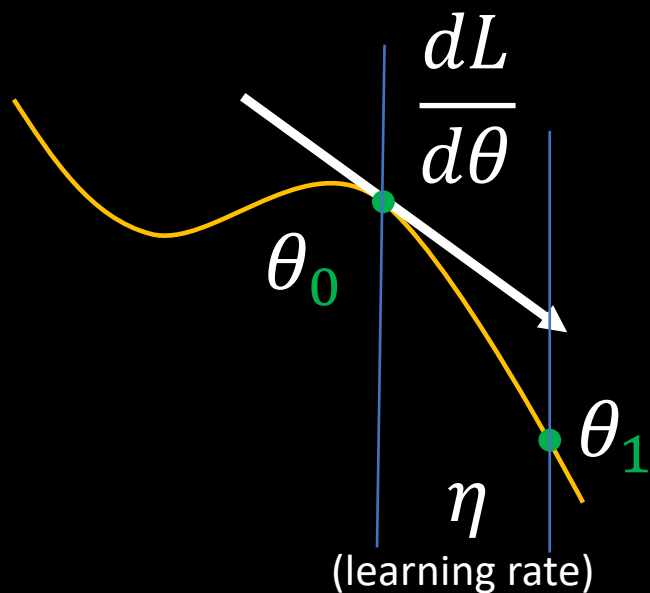
Moves in the direction of steepest descent (negative gradient) to minimize the objective function.

Steps:

1. Start with initial parameter values
2. Calculate the gradient of the objective function
3. Update parameters in the opposite direction of the gradient
4. Repeat steps 2-3 until convergence



Gradient Descent



$$\theta = \operatorname{argmin}(\sum_n (y_n - \hat{y})^2) =$$

$$\operatorname{argmin}(\sum_n (y_n - wx_n - b)^2)$$

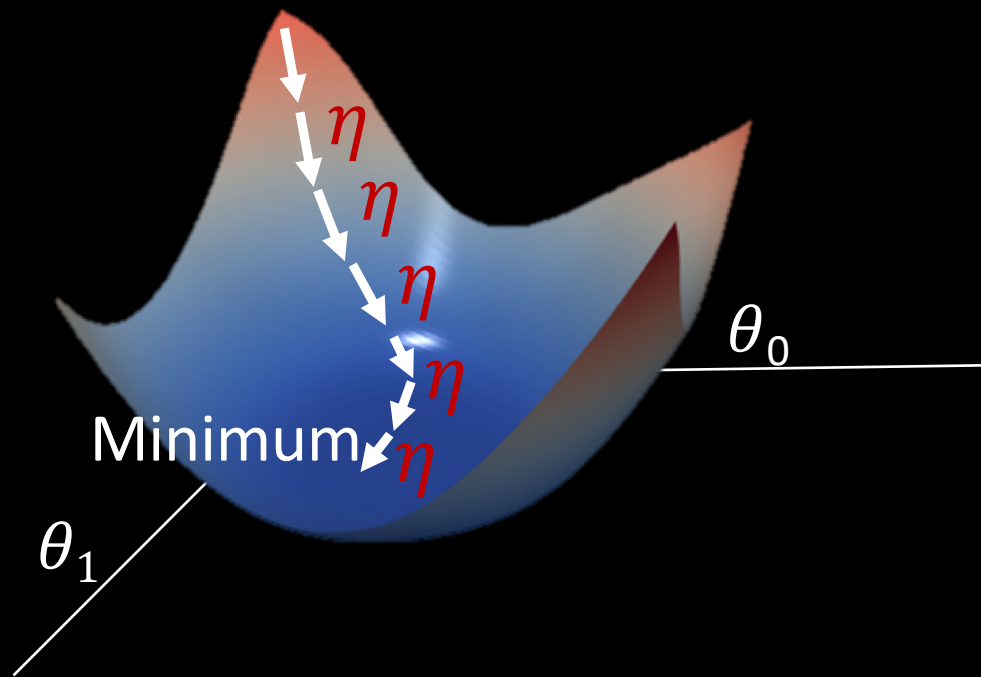
$$L = \sum_n (y_n - wx_n - b)^2$$

Loss function

$$\frac{dL}{dw} = \frac{d}{dw} \sum_n (y_n - wx_n - b)^2$$

Gradient of L
to parameters

$$\frac{dL}{db} = \frac{d}{db} \sum_n (y_n - wx_n - b)^2$$



$$\frac{dL}{dw} = \frac{dL}{dw} \sum_n (y_n - wx_n - b)^2 = - \sum_n 2x_n (y_n - wx_n - b)$$

$$\frac{dL}{db} = \frac{dL}{db} \sum_n (y_n - wx_n - b)^2 = - \sum_n 2(y_n - wx_n - b)$$

Remember chain rules?

$$F'(x) = f'(g(x)) g'(x)$$

```

import numpy as np

# Sample data
X = np.array([1, 2, 3, 4, 5])
y = np.array([2, 4, 6, 8, 10])

# Two parameter sets
w1, b1 = 1.5, 1.0
w2, b2 = 1.5, 0.5

# Function to calculate predictions
def predict(X, w, b):
    return w * X + b

# Function to calculate Mean Squared Error (MSE) loss
def mse_loss(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)

# Function to calculate gradient of loss with respect to b
def gradient_b(y_true, y_pred):
    return -2 * np.mean(y_true - y_pred)

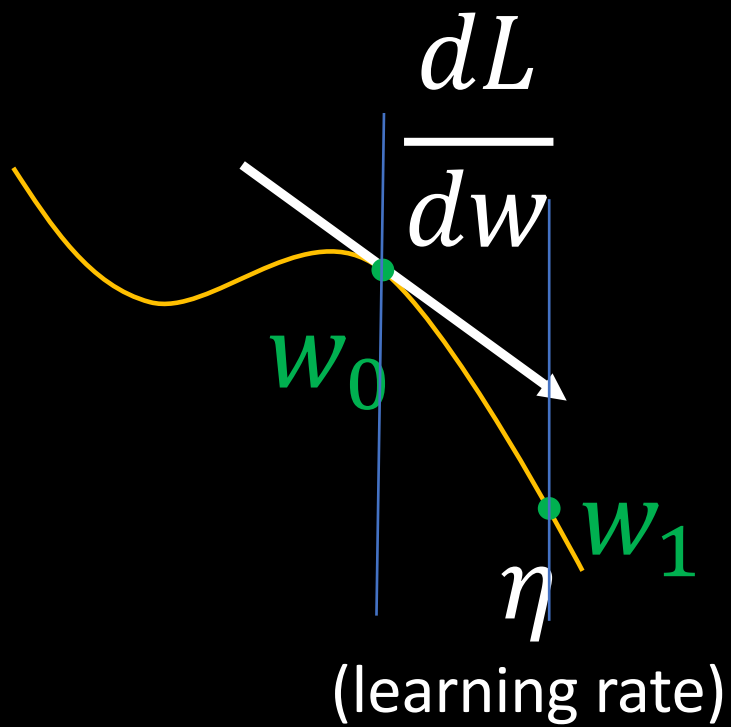
# Calculate predictions, loss, and gradient for both parameter sets
y_pred1 = predict(X, w1, b1)
loss1 = mse_loss(y, y_pred1)
grad_b1 = gradient_b(y, y_pred1)

y_pred2 = predict(X, w2, b2)
loss2 = mse_loss(y, y_pred2)
grad_b2 = gradient_b(y, y_pred2)

```

Parameter set 1: $w = 1.5$, $b = 1.0$
 Loss: 0.7500
 Gradient dL/db : -1.0000

Parameter set 2: $w = 1.5$, $b = 0.5$
 Loss: 1.5000
 Gradient dL/db : -2.0000



$$w_1 = w_0 - \frac{dL}{dw} \eta$$

$$b_1 = b_0 - \frac{dL}{db} \eta$$

$$(w_0, b_0) \rightarrow (w_1, b_1)$$

```
import numpy as np
X = np.array([1, 2, 3, 4, 5])
y = np.array([2, 4, 6, 8, 10])

# Initial parameters
w, b = 1.5, 1.0

# Function to calculate predictions
def predict(X, w, b):
    return w * X + b

def mse_loss(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)

def gradient_w(X, y_true, y_pred):
    return -2 * np.mean(X * (y_true - y_pred))

def gradient_b(y_true, y_pred):
    return -2 * np.mean(y_true - y_pred)
```

```
# Calculate initial predictions, loss, and gradients
```

```
y_pred = predict(X, w, b)
initial_loss = mse_loss(y, y_pred)
grad_w = gradient_w(X, y, y_pred)
grad_b = gradient_b(y, y_pred)
```

```
# Perform one step of gradient descent
```

```
learning_rate = 0.01
w_new = w - learning_rate * grad_w
b_new = b - learning_rate * grad_b
```

```
# Calculate new predictions and loss
```

```
y_pred_new = predict(X, w_new, b_new)
new_loss = mse_loss(y, y_pred_new)
```

```
print(f"Initial parameters: w = {w}, b = {b}")
print(f"Initial loss: {initial_loss:.4f}")
print(f"Gradient dL/dw: {grad_w:.4f}")
print(f"Gradient dL/db: {grad_b:.4f}")

print(f"\nAfter one gradient descent step:")
print(f"New parameters: w = {w_new:.4f}, b = {b_new:.4f}")
print(f"New loss: {new_loss:.4f}")

print(f"\nLoss reduction: {improvement:.2f}%")
```

Initial parameters: w = 1.5, b = 1.0

Initial loss: 0.7500

Gradient dL/dw: -5.0000

Gradient dL/db: -1.0000

After one gradient descent step:

New parameters: w = 1.5500, b = 1.0100

New loss: 0.5206

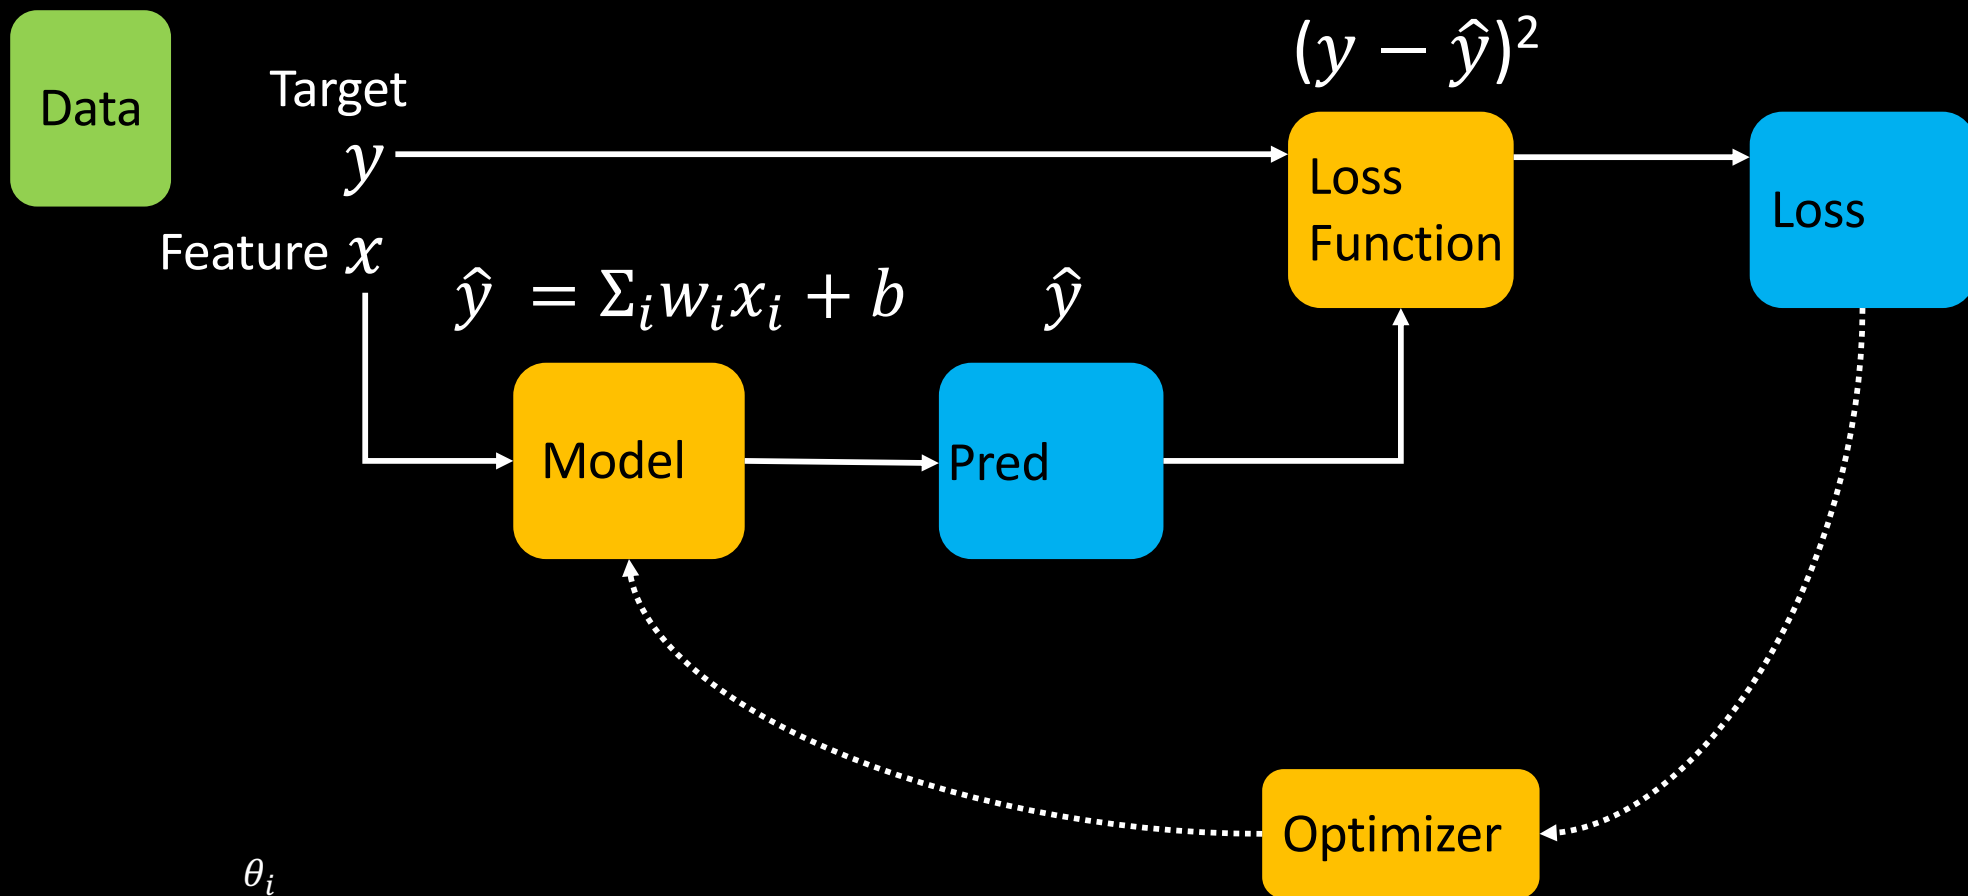
Gradient Direction:

- Gradient ∇L points towards the steepest increase in loss
- We move in the opposite direction ($-\nabla L$) to minimize loss
- For each parameter θ : $\theta_{new} = \theta_{old} - \eta * \partial L / \partial \theta$

Learning Rate (η):

- Controls the step size in each iteration
- Too large: May overshoot the minimum, causing divergence
- Too small: Slow convergence, may get stuck in local minima

Flow Chart for Linear Regression



Graph of (Multi-variable)Linear Regression

